

3. SPARK STREAMING

Apache Spark - December 2021

1. Introduction to Spark Streaming
2. Structured Streaming
3. Windowing
4. Kafka Connectors

1. Introduction to Spark Streaming



2. Structured Streaming

3. Windowing

4. Kafka Connectors

Motivation for Real-Time Processing

- **Data is being created at unprecedented rates**
 - Exponential data growth from mobile, web, social
 - Connected devices: 9B in 2012 to 50B by 2020
 - Over 1 trillion sensors by 2020
 - Datacenter IP traffic growing at CAGR of 25%
- **How can we harness it data in real time?**
 - Value can quickly degrade → Capture Value immediately
 - From reactive analysis to direct operational impact
 - Unlocks new competitive languages
 - Requires a completely new approach...

Uses Cases Across Industries

Credit



Identify fraudulent transactions as soon as they occur.

Transportation



Dynamic Re-routing Of traffic or Vehicle Fleet.

Retail



- Dynamic Inventory Management
- Real-time In-store Offers and recommendations

Consumer Intel Mobile



Optimize user engagement based on user's current behavior.

Healthcare



Continuously monitor patient vital stats and proactively identify at-risk patients.

Manufacturing



- Identify equipment failures and react instantly
- Perform Proactive maintenance.

Surveillance



Identify threats and intrusions In real-time

Digital Advertising & Marketing



Optimize and personalize content based on real-time information.

Introduction

- Spark Streaming provides a scalable, fault tolerant, efficient way of processing streaming data using Spark's simple programming model
- It converts streaming data into “**micro batches**”, which enable Spark's batch programming model to be applied in Streaming use cases
- This unified programming model **makes it easy to combine batch and interactive** data processing with streaming

1. Introduction to Spark Streaming
2. Structured Streaming 
3. Windowing
4. Kafka Connectors
5. Caching / Persistence
6. Fault Tolerance & Reliability

Structured Streaming

DStreams

Based on RDDs
Micro-batching
Non-Structured
Missing event time,
watermarking, late data, ...

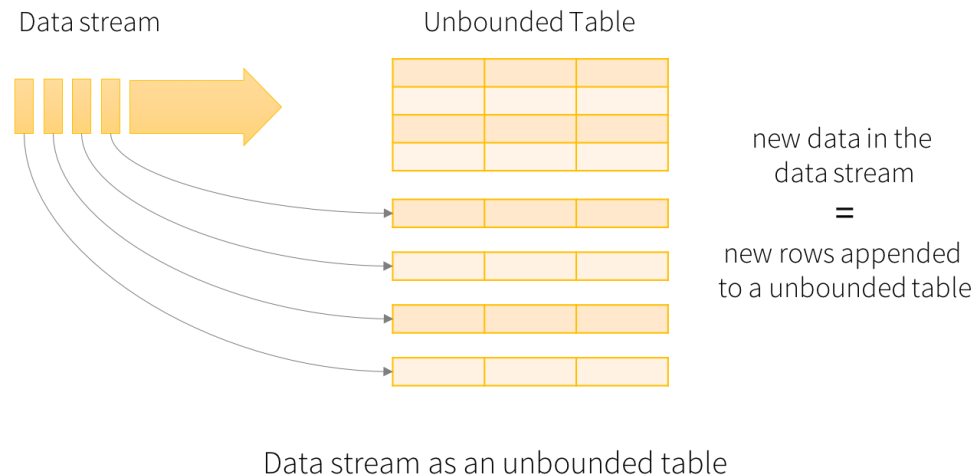


Structured Streaming

Spark v2.0+
Dataframes/Datasets
Catalyst Optimizer
Watermarking
Output modes: complete,
append, update

Structured Streaming

- Structured Streaming is a **scalable and fault-tolerant** stream processing engine built on the Spark SQL engine
- You can express your streaming computation the same way you would express a batch computation on static data



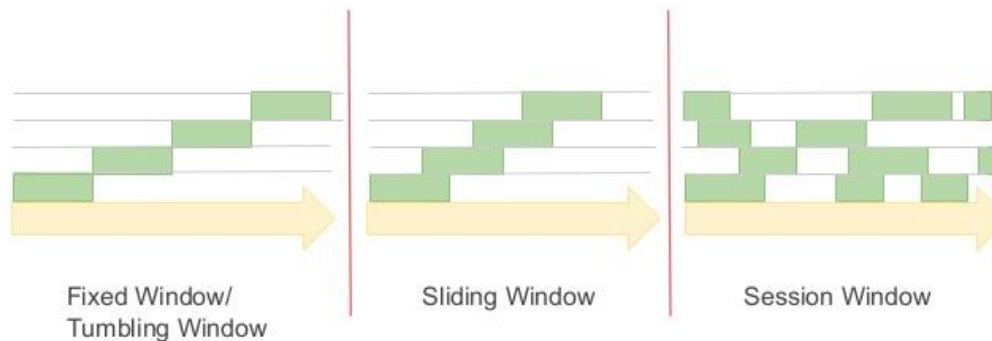
1. Introduction to Spark Streaming
2. Structured Streaming
3. Windowing
4. Kafka Connectors
5. Caching / Persistence
6. Fault Tolerance & Reliability



Windowing

- Windowing is the ability to perform some set-based computation (aggregation) or other operations over subsets of events that fall within some period of time

Streaming Concepts - Windows



```
df.withWatermark("timestampColumn", "5 hours")  
  .groupBy(window("timestampColumn", "1 minute"))  
  .count()
```

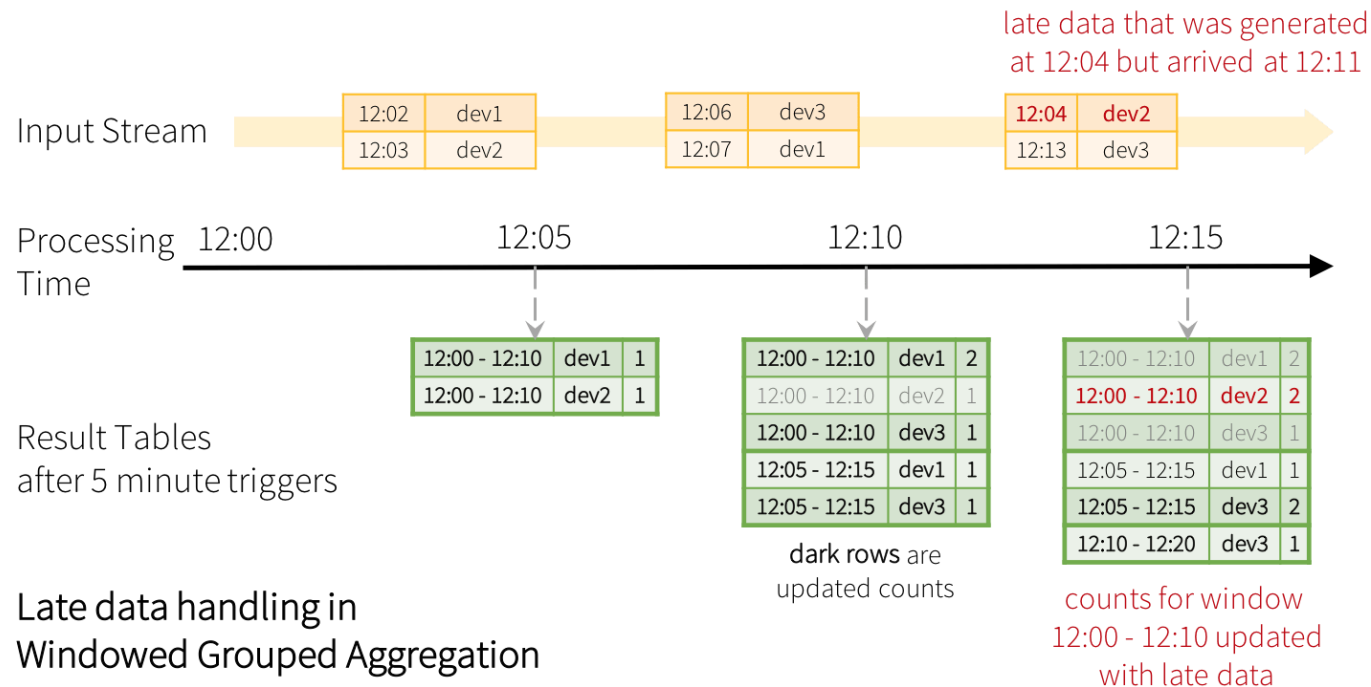
Same in streaming & batch

How to group
data by time



Watermarking

- Watermarking is a moving threshold in event-time that trails behind the maximum event-time seen by the query in the processed data



Output Modes

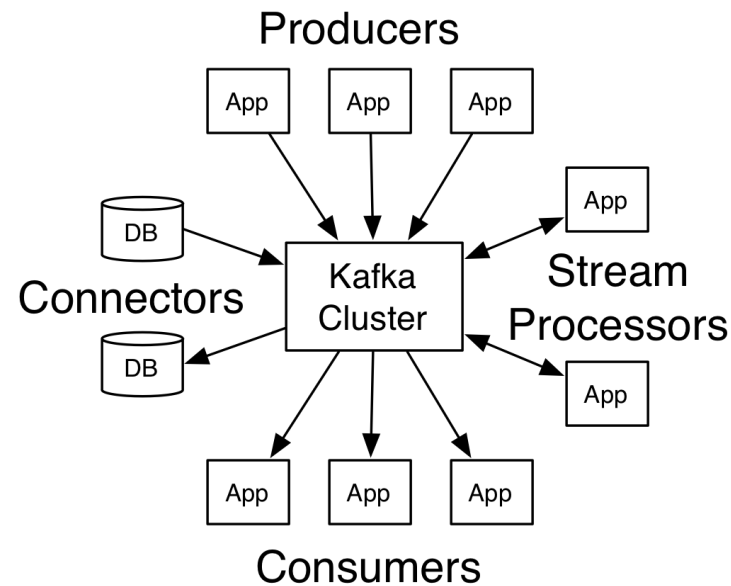
- **Complete Mode** The entire updated Result Table will be written to the external storage.
- **Append Mode** Only new rows appended in the Result Table since the last trigger will be written to the external storage.
- **Update Mode** Only rows that were updated in the Result Table since the last trigger will be written to the external storage (available since Spark 2.1.1)
- Not all output modes are feasible with all queries; check [here](#)

1. Introduction to Spark Streaming
2. Structured Streaming
3. Windowing
4. Kafka Connectors
5. Caching / Persistence
6. Fault Tolerance & Reliability



Apache Kafka Architecture

- Apache Kafka is a distributed streaming platform with three key capabilities
 - Publish and subscribe to streams of records
 - Store streams of records in a fault-tolerant durable way
 - Process streams of records as they occur



Kafka Connectors

- Structured Streaming integration for Kafka 0.10 to read data from and write data to Kafka.

```
// Subscribe to 1 topic  
val df = spark  
  .readStream    //read for batch queries  
  .format("kafka")  
  .option("kafka.bootstrap.servers", "host1:port1,host2:port2")  
  .option("subscribe", "topic1")  
  .load()  
  
df.selectExpr("CAST(key AS STRING)", "CAST(value AS STRING)").as[(String, String)]
```


Kafka Connectors

```
// Write key-value data from a DataFrame to a specific Kafka topic specified in
an option
val ds = df
  .selectExpr("CAST(key AS STRING)", "CAST(value AS STRING)")
  .writeStream
  .format("kafka")
  .option("kafka.bootstrap.servers", "host1:port1,host2:port2")
  .option("topic", "topic1")
  .start()

// Write key-value data from a DataFrame to Kafka using a topic specified in the
data
val ds = df
  .selectExpr("topic", "CAST(key AS STRING)", "CAST(value AS STRING)")
  .writeStream
  .format("kafka")
  .option("kafka.bootstrap.servers", "host1:port1,host2:port2")
  .start()
```

Hands-on

- Open “03.Spark_Streaming_Kafka.ipynb” in Google Colab:
 - Execute example 1
 - Try Exercises 1, 2 and 3

